

# A Tunable Compression Framework for Bitmap Indices

**Gheorghi Guzun, Guadalupe Canahuate**

The University of Iowa

**David Chiu**

Washington State University

**Jason Sawin**

University of St. Thomas

# Overview

- Big Data and bitmap indexing
- Background on bitmap compression
  - WAH, PLWAH, EWAH, COMPAX.
- Variable Aligned Length (VAL) Framework
  - Bitmap Compression – Segment Length Selection
  - Query Execution – between bitmaps encoded with different lengths
  - User defined parameter to control trade-offs
- Evaluation
- Conclusions and Future work


# Big Data and Bitmap Indices

- Big Data enables deeper understanding of things but comes at a cost
  - Disk space, memory and bandwidth needs grow very fast
  - When used properly, bitmap indices improve both compression and query times
- A Bitmap Index is a binary vector that represents which elements fall into a category or have certain property
  - Used in data warehousing applications with large amounts of data and ad hoc queries, but a low level of concurrent DML transactions.
- Advantages:
  - Leverage fast bit-wise operations to resolve queries
  - Highly compressible (hybrid run-length encoding)



# Bitmap Example

Tuple	Name	Age	Salary	City	Favorite Drink	...
<i>t1</i>	Sara	42	65,000	Beaverton	Ninkasi	
<i>t2</i>	Julie	50	130,000	West Linn	Bridgeport	
<i>t3</i>	Tom	21	25,000	Portland	Boneyard	
...						



*Discretize (binning) and bit assignment*

Name				Age			
Tuple	Sara?	Julie?	Tom?	Age <= 25	25 < Age < 50	50 <= Age	...
<i>t1</i>	1	0	0	0	1	0	...
<i>t2</i>	0	1	0	0	1	0	
<i>t3</i>	0	0	1	1	0	0	
...							

# Bitmap Compression

- Goal:
  - Reduce bitmap size
  - Minimize overhead in query execution time
    - Execute query without explicit decompression
- Run-length encoders
  - Word-Aligned Hybrid Code (WAH) – widely adopted
    - Two types of words - simple
    - Word Aligned → Faster CPU time
    - EWAH, PLWAH, CONCISE, COMPAX are word-aligned based

133 Bits	1,20*0,4*1,78*0,30*1
31-bit groups	1,20*0,4*1,6*0    62*0    10*0,21*1    9*1
groups in hex	400003C0 00000000 00000000 001FFFFFF 000001FF
WAH (hex)	400003C0 80000002 001FFFFFF 000001FF

Literal Word    Fill Word

# WAH Compression – Weak points

- Often uses too many bits for encoding runs
  - 30 bits for  $W=32$  and 62 bits for  $W=64$
  - For  $W=64$ , longest compressible run in one word is  $63 \cdot 2^{62}$
  - PLWAH, CONCISE, COMPAX try to improve by exploiting this waste of bits
- Not all run-lengths are equal
  - Most real-world datasets are skewed
  - Sorted bitmaps have even greater contrasts in run-lengths
- **Solution – Variable number of bits to compress runs**

# Variable Length Compressions

- Variable Length Compression (VLC)
  - Uses variable number of bits (variable segment length) for compression
  - Looser alignment requirements – slows down queries dramatically
- PWAH (Partitioned WAH)- WAH using shorter segment lengths (partitions)
  - Encodes entire dataset with one single segment length
  - Does not provide query algorithms between bitmaps encoded with different segment lengths

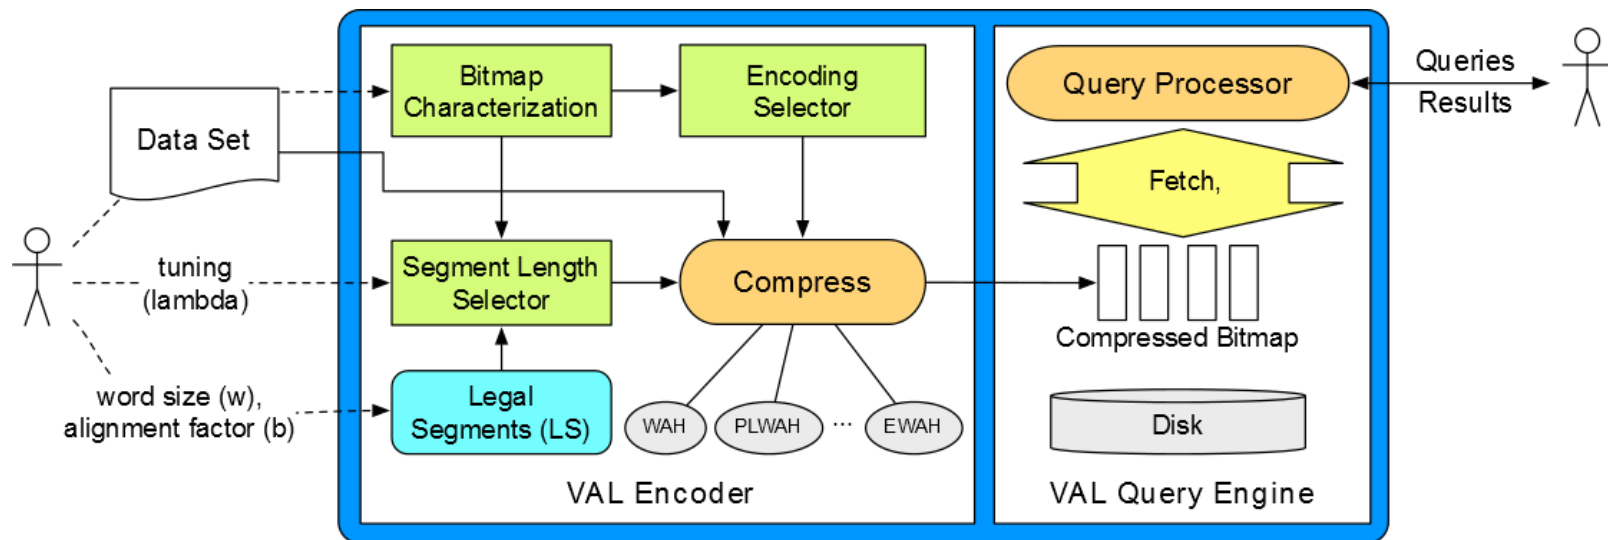
# VAL- Variable Aligned Length (Challenges)

- Improve Compression
  - Use variable number of bits (segments) to compress runs  $s < W$
- Keep the word encapsulation structure
  - Reading entire words streamlines memory accesses
- Efficient query algorithms
  - Enable queries between bitmaps encoded using a different “s”
  - Maintain alignment between segments



# The VAL Framework

- Integrates existing compression techniques – Uses VAL
  - Based on the observation that most existing word aligned based encodings use the essentially the same query algorithm

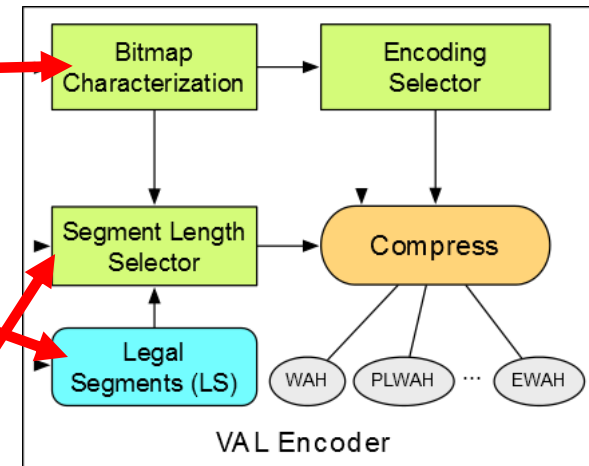


# Segment Length Selection

- Scan data
  - Collect bit-density, bitmap size, run-lengths
- Legal Segment Lengths
  - $LS = \{2^i \times (b - 1) | 0 \leq i \leq (\log_2(w) - \log_2(b))\}$
  - e.g. for  $W=64$ , LS can be  $\{15, 30, 60\}$


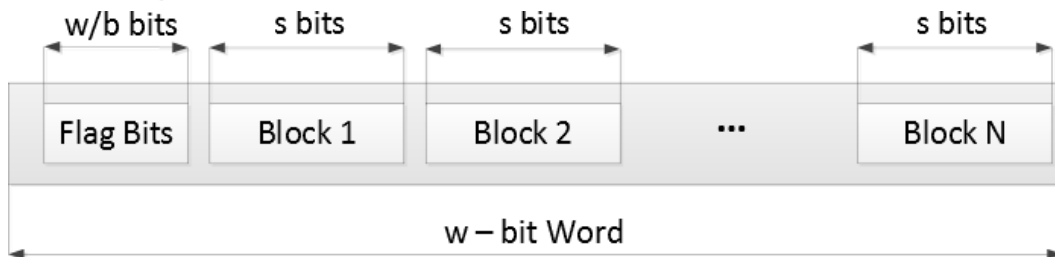
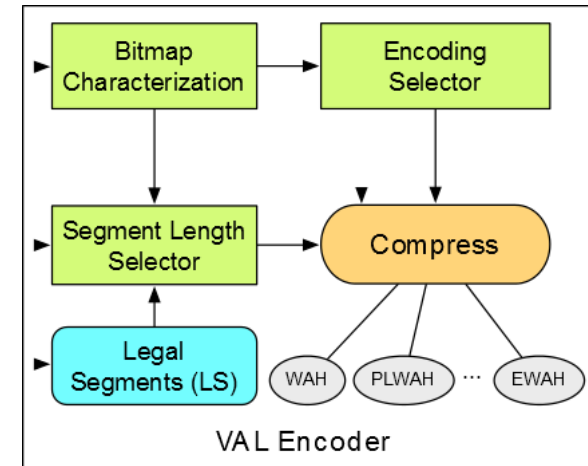
## Segment Length Selector

- $$\begin{cases} s_{c+i}, & \text{if } \frac{\text{size}(B, s_c) \times (1+\lambda)^{1+i+\lambda}}{i+1} \geq \text{size}(B, s_{c+i}) \\ s_c, & \text{Otherwise} \end{cases}$$
- $s_c = \arg \min_{s_i \in LS} \{\text{size}(B, s_i)\}$




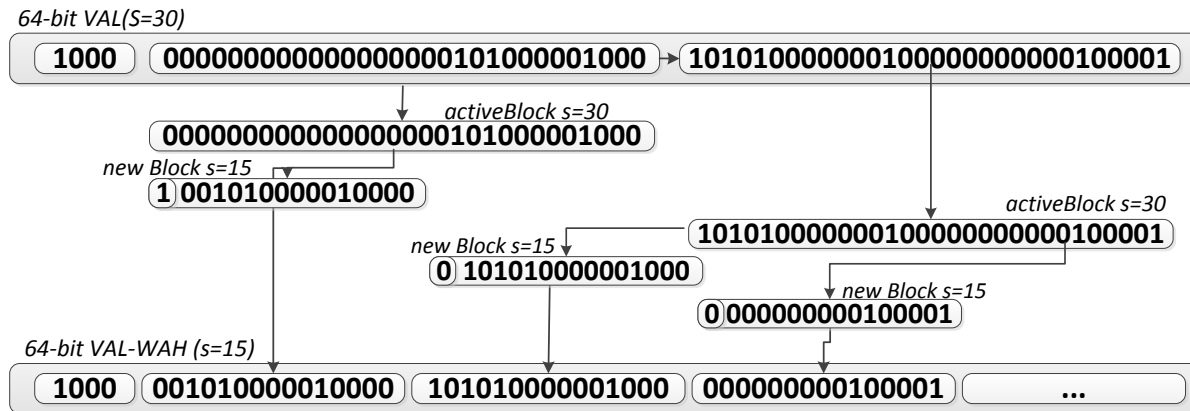
# Compression

- Compresses given selected method and segment length
- Encapsulates blocks into words
- Each word contains flag-bits for it's blocks
- Each bitmap has a header where the segment length is stored



# General Query Algorithm

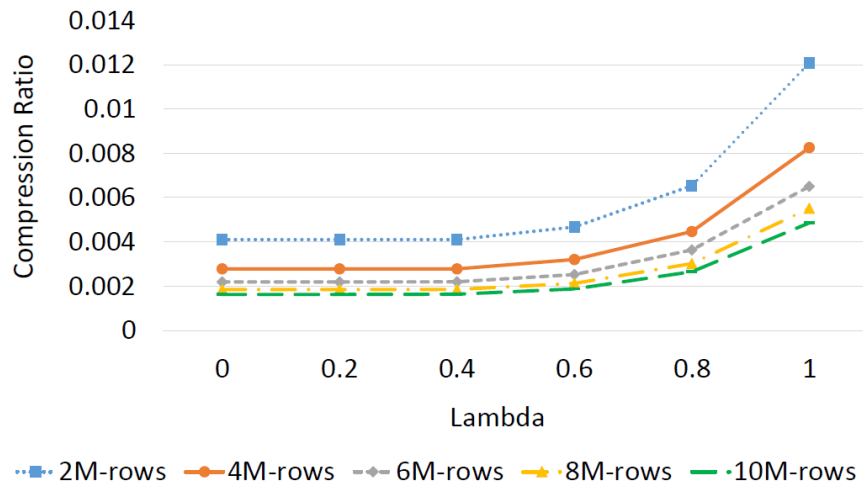
- One general logical operation algorithm for all encodings
  - Iterate the two bitmaps and operate block by block
  - Convert between segment lengths if necessary



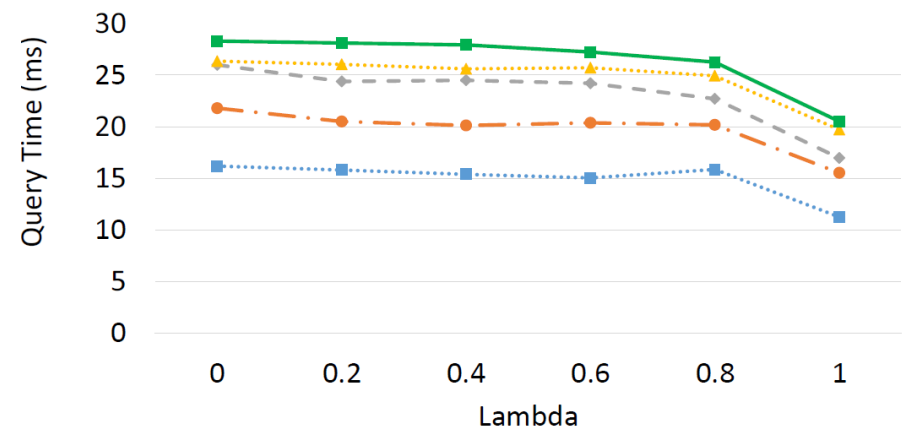
# Experimental Setup

- Implemented VAL-WAH as proof of concept
- Run on 64-bit Intel i7-2600 3,2 GHz 8MB cache and 8 GB RAM
- Synthetic datasets
  - Use zipf distribution with skew 0; 1;2.
  - Sorted and non-sorted. For sorted use gray-code sorting
- Real datasets
  - Kddcup, Berkeley Earth

# Effects of lambda

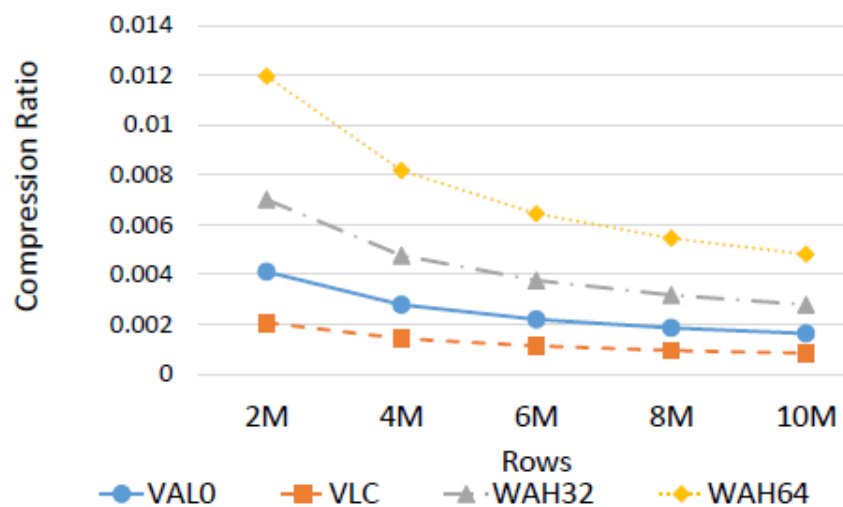


(a) Synthetic Data (zipf2) - Sorted: Compression Ratio




(b) Synthetic Data (zipf2) - Sorted: Query Time

# Segment alignment





(a) Synthetic Data (zipf2) - Sorted: Compression Ratio



(b) Synthetic Data (zipf2) - Sorted: Query Time

# Real Data


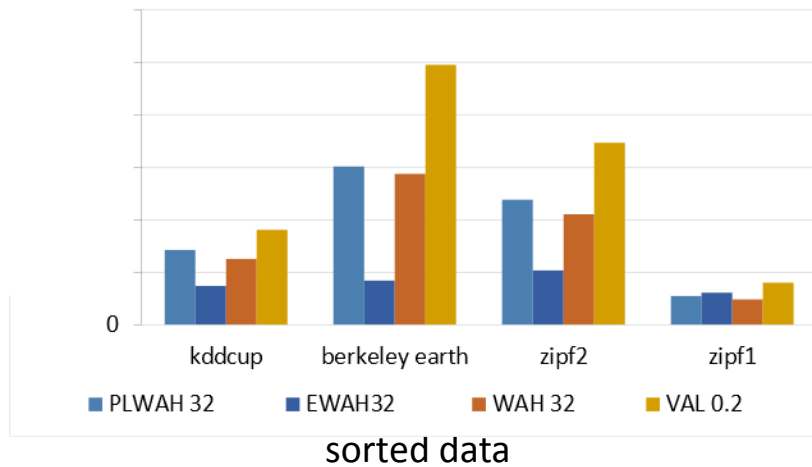




# Combined gain

- To help simplify discussion on trade-off:

- $gain = \frac{1}{H_m} = \frac{queryTimeRatio + compRatio}{2 \times queryTimeRatio \times compRatio}$



# Summary

- We developed the VAL framework
  - Uses variable segment lengths for better compression
  - Maintains alignment for faster query processing
  - Introduces a tuning parameter to adjust trade-offs
- Future work
  - Implement other word-aligned encodings within VAL
  - Develop algorithms for selecting encoding methods