

Incremental Quantization for Aging Data Streams

Fatih Altıparmak David Chiu Hakan Ferhatosmanoglu [†]
 Department of Computer Science and Engineering
 The Ohio State University
 {altiparm, chiud, hakan}@cse.ohio-state.edu

Abstract

A growing number of applications have become reliant or can benefit from monitoring data streams. Data streams are potentially unbounded in size, hence, Data Stream Management Systems generally maintain a “sliding window” containing the N most recent elements. In an environment where the number of stream sources can vary, the amount of storage available to hold the sliding window can reduce dramatically. However, it has already been noted that as data becomes older their relevance tends to diminish until they are ultimately discarded from the sliding window. Based on this assumption, we propose to “wound” older data elements by relaxing their storage requirements as an effort constantly free enough space to keep pace with accurate representation of incoming elements in a process that we call aging. We propose two incremental quantization techniques that enable aging in an efficient manner. We will show that, by relaxing storage utilization of the summary created by our quantizers, the older data elements are not rendered useless. In fact, we will show that their accuracy is only lessened by a sustainable amount.

1 Introduction

Data Stream Management Systems (DSMS) have gained significant interest in the database community due to a growing number of applications that deal with continuous data, e.g., environmental monitoring and network intrusion detection. The general challenge that comes coupled with the data stream model is storage management. Once an element from a data stream has been processed it is discarded or archived - it cannot be retrieved easily unless it is explicitly stored in memory, which typically is small relative to the size of the data streams [2]. Because data streams are potentially unbounded in size, storing the complete summary in memory is prohibitive. To overcome this issue,

DSMS's generally maintain a constant size “sliding window” containing the N most recent data stream elements. By restricting the stream application to finite bounds, it becomes feasible to store and access the most recent elements of streaming data efficiently, assuming that real-time applications mostly find recent data more relevant.

Consider multi-streaming models where M streams feed into a central DSMS server, which must maintain a sliding window over these streams. When M becomes large, system resources (e.g., CPU and storage) are consumed at a much more vigorous rate, becoming increasingly difficult to maintain Quality of Service constraints such as timeliness and accuracy. In order to alleviate CPU load under these strenuous conditions, recent development in DSMS design employs techniques such as accuracy adjustment and load shedding [4, 5, 13]. The burden on storage has also received attention. Recall that the sliding window must be constant size in order to fit in memory for efficient processing. Summarization techniques including histograms [6, 7, 10], sketching [8] and sampling [3] pave the way to data reduction.

We propose an approach to mitigate the stress on storage in multi-stream systems by adding the element of *age* to the sliding window. Aging data streams is not a novel concept as Jagadish et al. [10] first proposed to take into account the age of the elements. Their approach aims to provide fast approximate point queries by updating a histogram summary. Ours differ in the way that we offer an online compression technique to hold the actual values (not just statistics) of the sliding window with minimal loss of accuracy, thus enabling a wide array of data mining tasks including clustering and such as K -Nearest Neighbors queries. The intuition is established on the observation that stream applications carry natural tendencies to favor recent elements over older elements [12]. As data grows older, its relevance diminishes until it is ultimately discarded from the sliding window. Our technique “wounds” older data elements by relaxing their storage requirements. This regained space is then further utilized to capture higher accuracy for incoming elements.

We make several contributions in this paper. First,

[†]This research is partially supported by US National Science Foundation grants III-054713, CNS-0403342, and OCI-0619041.

we propose 2 novel online quantization techniques, *Incremental V-Optimal Quantization (IVQ)*, and *Distortion Based Quantization (DBQ)* that achieve efficient on-demand compression. Correspondingly, we will also show that not only do *IVQ* and *DBQ* find near-optimal solutions, i.e., minimal loss in quantized data representation. We find that they are also polynomially faster than current approaches to the optimal solution. Finally, we discuss the *Ladder Approach* for sensible storage distribution as a means to apply age to data elements.

The rest of this paper is organized as follows. In Section 2, we describe the role of quantization with regards to supporting aging on streams. We then propose *IVQ* in Section 3. Next, Section 4 discusses the Ladder Approach towards aging data stream elements, followed by a performance analysis of our approach in Section 5. Finally, we conclude and discuss future directions in Section 6. Due to space constraints, a full discussion of *DBQ* can be found in [1].

2 Background

Because quantization lies deep in our aging procedure, we first present a brief review of this technique. The goal of quantization is to constrain a given vector F of continuous values into B discrete intervals (with $B - 1$ boundaries). Each discrete interval can be uniquely identified by a label consisting of $\lceil \log_2(B) \rceil$ bits. Thus, the actual values in F are replaced with a $\lceil \log_2(B) \rceil$ bit sequence which indicates the the interval with its closest representation. Albeit a loss in accuracy of the actual elements, this technique has been proven useful in many disciplines including image processing and signal encodings. Moreover, as the number of intervals increases the information lost due to the quantization is reduced. We illustrate this procedure with a simple example.

$F =$	0	2	3	6	7	15	16	20
Interval	0	0	0	0	0	1	1	1

Using $b = 1$ bit for quantization, the domain, $D = [0 \dots 20]$, of the 8 given numbers is divided into $B = 2$ intervals, I_0 and I_1 , separated by $B - 1 = 1$ boundary. For sake of simplicity, we let the boundary $\psi = 10$. Then all given values $f \in F$ that are less than or equal to 10 are assigned to I_0 and the rest are assigned to I_1 . Notice that only a single bit is needed for mapping all values in F to the closest interval.

Statistics can also be applied per interval to increase the information denoted by the bit representation. Typically, the mean or median of all values within each interval is used for this purpose. In order to maintain an accurately quantized synopsis, now the goal is to minimize the sum of the square of the differences (SSE) between all values within

an interval and the statistic of the interval to which they are assigned. In this example, 3.6 and 17 are the means of I_0 and I_1 respectively, and $SSE_0 = 33.2$ and $SSE_1 = 14$. The problem becomes finding the interval boundaries such that they minimize $SSE_{total} = \sum SSE_k$.

Although our boundary was trivially assigned to 10 in this example, effective boundary identification is a non-trivial task on real number domains. For instance, if the boundary had been selected as $\psi = 5$, then the points 6 and 7 would have fallen into I_1 . Because 0, 2, and 3 are very close to each other, this would expectedly decrease SSE_0 to 4.67. However, it increases SSE_1 to 178.2, and $SSE_{total} = 182.78$. With regards to storage, assuming that it takes 4 bytes to represent integers and floating point numbers, without quantization the given values in F would allot 32 bytes. On the other hand, quantization observes the following attributes on storage: $4B$ bytes to store the B interval statistics, and bM bits for the M b -bit representations used for interval mapping. Quantizing F generates only 8 bytes + 8 bits = 9 bytes to hold the lossy representation of the given values. Considering the fact that we must sustain some loss in data accuracy, recall that the main goal of quantization is to use minimal space to maintain a maximally accurate summary. It should now be clear that optimally minimizing SSE_{total} translates to this solution.

2.1 Enabling Aging

In our system aging refers to the process of wounding older data elements by relaxing their storage and then using this storage to capture the details of the synopsis belonging to the younger incoming elements. Hence, the design of the quantizer should be made appropriate not only to support such an operation, but to support it efficiently. We summarize the requirements that should be met by the quantizer as follows: (1) Efficient Aging — the final quantization function should be built from an incremental manner in such a way that wounding a bit from the synopsis takes minimal time. (2) Preserving High Accuracy — boundaries should be selected in such a way that SSE_{total} is minimized.

While the first requirement will be discussed in the next section, the second requirement refers to solving the V-Optimal Histogram Problem, which states, given a vector of M numbers and the number of desired intervals, B , find the $B - 1$ boundaries such that SSE_{total} is minimized. While optimal solutions to this problem are polynomial in nature [9, 10, 6], there exists a notion that, when M is large, approximate solutions become more practical [10, 7, 11]. Unfortunately, none of these solutions fit our first requirement of an incremental quantizer construction. We show the need for an incremental design through an example with quantizing multi-stream data.

In Figure 1, assume that our system handles $M = 10$

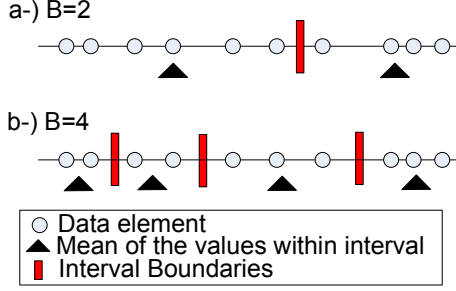


Figure 1. Challenge of the 1st Criterion

streams, and we let each point denote a single element arriving from a data stream. Suppose that these elements are currently represented in our sliding window with a $b = 2$ bit quantization, as depicted in Figure 1b. After some time, our system recommends that these data are now “older” and should be compressed to make space for incoming elements. In order to relax a bit from the 2-bit representation, our system must store a corresponding 1-bit representation as shown in Figure 1a. Without a directly accessible 1-bit representation, it would become necessary to again quantize the 2-bit data using one less bit — an undesirable and seemingly redundant process. Furthermore, we would also expect too much error to be introduced by the re-quantization of already quantized (lossy) data. Another approach is to save both representations independently. Although it would guarantee a statically accessible representations, the algorithm would have to be run b times to obtain each. Moreover, its increased storage requirements for each set of data makes this option even more so unlikely for data stream systems. It is clear that a quantization algorithm is needed such that the boundaries of the quantization for b bits are preserved during the quantization for $b + 1$ bits. That is, a b bit quantization should be constructed iteratively such that new boundaries are added on top of existing boundaries. We note that the current solutions to the V-Optimal Histogram Problem are not incremental, and therefore unable to apply aging efficiently.

Incremental quantization techniques do, however, exist. Most notably, equal-length and equal-depth quantization are heuristics that naturally observe incremental behavior [10]. Their interval boundary selection criteria are resident in their names: equal-length iteratively partitions an interval into 2 equal halves while equal-depth divides intervals into halves containing an equal number of elements. While both techniques are mechanically efficient and incremental, the trivial boundary selection criteria may produce poor data accuracy on practical distributions. In light of this issue we propose a novel V-Optimal Histogram heuristic that identifies near optimal *incremental* intervals while achieving the same linear time complexity as equal-length and

equal-depth.

3 Iterative V-Optimal Quantization

The proposed method, Iterative V-Optimal Quantization (*IVQ*), exploits a special case in the solution for V-Optimal Histogram. The major issue with directly employing the exact optimal solution [6] is its polynomial time complexity $O(M^2)$. However, we prove that for the special case of $B = 2$, finding an optimal boundary, ψ , for 2 intervals can be done in $O(M)$ -time on a sorted set of M numbers. The premise of *IVQ* starts with a 0-bit summary, then iteratively using this $O(M)$ -time subroutine, we add 1 bit to our summary at a time until the number of total bits is satisfied. The incremental construction of the quantized summary allows us the structure to remove bits efficiently. But first, we prove that when $B = 2$, the optimal solution can be found in $O(M)$ -time.

We call this special case of $B = 2$ the *V-Optimal_{B2}* Problem. Given a sorted vector F of length M , we can state this problem as follows: find i such that

$$\sum_{k=1}^i (F[k] - \text{Avg}(F[1 \dots k]))^2 + \sum_{k=i+1}^n (F[k] - \text{avg}(F[i+1 \dots n]))^2 \quad (1)$$

is minimized and $\psi = \frac{F[i] + F[i+1]}{2}$. We can rewrite equation 1 to obtain

$$\sum_{k=1}^n (F[k])^2 - (i * \text{avg}(F[1 \dots i])^2 + (n - i) * \text{avg}(F[i+1 \dots n])^2). \quad (2)$$

Now we see that minimizing equation 2 is equivalent to maximizing

$$i * \text{avg}(F[1 \dots i])^2 + (n - i) * \text{avg}(F[i+1 \dots n])^2. \quad (3)$$

Note that if we define an array P , as in [10], of length M with $P[j] = \sum_{1 \leq k \leq j} F[k]$ then we can rewrite equation 3 as

$$\frac{P[i]^2}{i} + \frac{(P[M] - P[i])^2}{M - i}. \quad (4)$$

After reformulation, we can see that finding the i that maximizes this equation is an $O(M)$ -time routine.

The *IVQ* algorithm, shown in 1, is summarized as follows. First, we divide the quantization process into b incremental steps, where b again denotes the desired number of bits used in our final representation. Assume that for each step, we use only 1 bit per data element for a given set of numbers. Recall that using 1 bit results in $B = 2$ intervals. This implies that at each step j , we need only to solve the *V-Optimal_{B2}* Problem to find the boundary, ψ_j . For the numbers in F on the left of the boundary, i.e., $< \psi$, the corresponding bit is assigned 0 and the rest 1.

Algorithm 1 The *IVQ* Algorithm

```

1   $b$  // number of bits to be used in final quantization
2   $psi[1 \dots B]$  // boundaries
3   $F[1 \dots M]$  // input vector
4   $list[] \leftarrow \text{sort}(F)$ 
5
6   $\text{IVQ}(list[], \text{start}, \text{end}, j, \text{psi\_index})$ 
7      if  $j \leq b$  then
8          {
9              //call solution for v-optimal-b2 to get boundary
10              $psi[\text{psi\_index}] \leftarrow \text{v-optimal-b2}(list, \text{start}, \text{end})$ 
11
12             //call recursively on both sides of lists
13
14             //first half
15              $\text{IVQ}(list, \text{start}, \text{psi}[\text{psi\_index}], j+1, 2*\text{psi\_index})$ 
16
17             //second half
18              $\text{IVQ}(list, \text{psi}[\text{psi\_index}] + 1, \text{end}, j+1, 2*\text{psi\_index}+1)$ 
19         }
```

To illustrate, when $j = 1$, we add the first bit to the synopsis. In do this, we first employ the above solution for V-Optimal_B2 Problem and find ψ_1 and the statistics of the two intervals. Then when $j = 2$, we add the second bit by apply our algorithm recursively to find ψ for each half. This process is continued until we reach $j = b$ whereby all $2^b - 1$ ψ 's (and $B = 2^b$ interval statistics) are computed. The complexity of *IVQ* is straightforward. For a list of size M , we run the $O(M)$ -time subroutine for *V-Optimal_B2*, then the list is split into 2 halves, and the algorithm is applied recursively on both halves. The number of splits is directly dependent on b . Therefore, cost of this algorithm is $O(bM)$, while keeping in mind that b is typically constant. The procedure is illustrated for $b = 3$ in Figure 2.

Recall that the original need for an incremental quantizer involved Requirement (1): efficient aging. Returning to Figure 2, when adding bits to the quantization summary, notice that the boundaries of the intervals from the previous iteration are preserved and reused. To visualize the process of aging, assume that we have a 3-bit representation of the data elements belonging to a particular time unit and we want to remove one bit from this representation. Because our summary was built in an incremental manner, we can simply remove the least significant bit from each stream in order to obtain the exact representation for the 2-bit summary. The removal of this bit only costs $O(1)$ -time. To support this low cost, we must maintain all statistics $Avg_{1,1}, Avg_{1,2}, Avg_{2,1}, \dots, Avg_{3,8}$, to map to the intervals. This implies saving $\sum_{i=1}^k 2^i$ means for some k -bit representation. The cost of this storage is amortized for a massive amount of streams (i.e., when M is large), but it should be noted that while deleting the k^{th} bit, we also can delete all averages belonging to the quantizer of that bit: $Avg_{k,1}, \dots, Avg_{k,2^k}$, thus regaining that space. As for achieving Requirement (2): accuracy, we address *IVQ*'s efficacy in data representation as compared to the optimal solution in Sec-

tion 5.

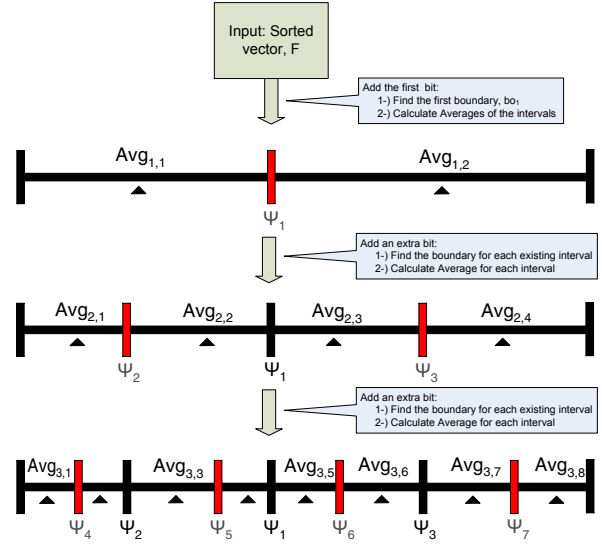


Figure 2. Example Run of *IVQ* for $b = 3$

Due to space constraints, we provide an offline description of our second incremental quantization method, *DBQ*, in [1]. However, we leave this technique in the performance evaluation section for completeness.

4 Overall Approach for Aging Data Streams

The concept of using *aging* on data streams to address the storage issue is born from the observation that newer streaming data is typically more relevant than the older [12]. This naturally corresponds to the idea of capturing more detail on the younger elements by sacrificing some detail from older elements. Our overall approach of supporting this is through storage manipulation: to utilize more storage for the data belonging to newer time units and less storage for older data, and devise a technique where more storage translates to higher accuracy.

We illustrate with an example. Assuming that the length of the sliding window is 5 time units, and additionally, the available memory for each stream is 15 units of storage. A simple way of distributing storage units would be to simply assign them uniformly across the sliding window: 3 bits per time unit. However, the interest of the queries is not quite as uniform. If most queries involve the first few time units, then we could sustain some depreciation of older elements. The storage can alternatively be assigned to the time units based on their order of the query interest. For instance, the youngest time unit obtains 5 units of storage, the second youngest gets 4 units, and so on.

We call this storage distribution method the *Ladder Approach*. At time $t + 1$, an element from each stream arrives

at our system. Focusing on a single stream, under the ladder model, being these youngest elements should receive 5 units of storage. The question now is where to obtain the needed space. We salvage 1 unit of storage from the element belonging to time $t - 4$ because it is now discarded from the sliding window. The data belonging at time t is no longer the youngest element in the synopsis, so it “steps down” the ladder and relaxes 1 bit of storage. The same holds for elements of $t - 1, t - 2$, and $t - 3$. The outcome of this aging operation are that there remains 5 units of available storage for the new element at time $t + 1$.

In general, the length of the sliding window, N , is much longer than the one given in our example. Assume that the total amount of storage available per stream is $\lceil(\omega + 1)/2\rceil * N$ units, where ω units of storage is presumably enough to capture nearly all details of the data belonging to a time unit after quantization. Then the synopsis can be maintained in such a way that the ladder contains ω steps. Based on the above assumptions, the time units can be assigned to the steps of the ladder uniformly. So, in the case of N is a dividend of ω , the number of time units using some u units of storage is equal to that using $u + 1$ units of storage. Otherwise, the N can be selected as the closest smaller multiplicand of ω . The starting time unit of the step s , for elements of which we are using s units of storage, can be calculated by $(t - N + 1) + (s - 1) * N/\omega$. As time passes from t to $t + 1$, the oldest time units in each step of the ladder are shed one unit of storage, which corresponds to deleting the synopsis itself for the oldest element. The ladder is then shifted towards the future by the aging operation discussed above. This is summarized in Figure 3.




Figure 3. The Ladder Approach for Storage Distribution

It should now be clear as to why an incremental quantization technique such as those presented in the previous sections are suitable underlying structures for enabling our aging approach. We have described how *IVQ* is purposely built in such a way that its incremental data structure supports a constant time routine to age the entire sliding window.

Focusing on our M stream system, our sliding window can be implemented by using nested linked lists. The outer list is a size N list where each node contains the summary of a specific time unit. This includes the boundary statistics and the bit mappings. The mappings itself is another list where each node contains M bits at some j th position. Aging on *IVQ* simply involves traversing each outer node and removing its first inner node (the most significant bit).

5 Experimental Results

In this section, we will compare the performance of all previously discussed quantization techniques with the optimal solution for 4 datasets of 10000 data points. The datasets follow the exponential, loglogistic, gaussian, and uniform distributions.

In Figure 4, the x-axis denotes the number of used bits, and the y-axis denotes the Mean Squared Error (MSE), which is simply the ratio SSE_{total}/M . The values closer to 0 signifies better quantization in terms of data accuracy.

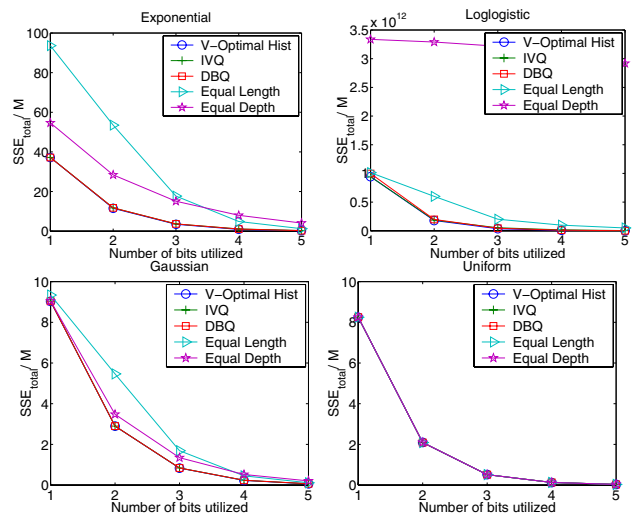


Figure 4. Comparison of different quantization techniques with respect to SSE_{total}/M

We compared the exact solution [10] for the V-Optimal Histogram Problem with the incremental techniques equal-depth, equal-length, *IVQ*, and *DBQ*. As can be seen in all graphs shown in Figure 4, *IVQ* and *DBQ* lead to smaller error ratios than equal-depth and equal-length. While the plots belonging to *IVQ* and *DBQ* ostensibly overlay that of the optimal solution in all 4 distributions, marginal errors do exist between these two techniques and the optimal. The major difference is that *IVQ* and *DBQ* are incremental, and fit our purpose for supporting aging, while V-Optimal does not (it would be too inefficient in that the wounding

process cannot be done efficiently).

As can be seen from the graphs in Figure 4, for 5 bits, the MSE for IVQ and DBQ techniques are very close to 0 for all datasets. The conclusion we can derive from these graphs is that a 5-bit quantization scheme seems to be enough to represent a given vector of size $M = 10000$ with these techniques with minimal noise.

In light of the above claim, we examined the effect of increasing the number of elements in these 4 datasets on the performance of the IVQ technique. Due to space restrictions, the remaining experimental results are shown only for IVQ . From each distribution we obtain a sample of 4 different sizes, $\{5000, 10000, 50000, \text{ and } 100000\}$. Table 1 shows the maximum ratio between $SSE_{total}/M * Variance(Data)$ for these datasets by using ≤ 6 bits for quantization. We use this normalization, $SSE_{total}/M * Variance(Data)$, in aim of comparing quantization techniques such that they can be qualified independently apart the actual data set. The concept is that SSE_{total} for 0-bit representation is equal to M times the variance of the data. By measuring the ratio $SSE_{total}/M * Variance(Data)$, we can now claim that the values closer to 0 signifies better quantization in terms of data accuracy, and those closer to 1 corresponds to noisier representation. As shown in Table 1, the maximum of this ratio increases only slightly when we increase the size of the datasets. This supports our claim that a 5-bit quantizer is adequate such that SSE_{total} is minimized.

Size of Datasets	1-Bit	2-Bit	3-Bit	4-Bit	5-Bit	6-Bit
5000	.4970	.1754	.0568	.0169	.0039	.0009
10000	.4922	.1689	.0503	.0142	.0034	.0009
50000	.4970	.1750	.0550	.0163	.0045	.0012
100000	.5002	.1766	.0545	.0160	.0044	.0011

Table 1. Max ratio between $SSE_{total}/M * Variance(Data)$ for datasets from 4 distributions

6 Conclusion and Future Work

As an effort to address storage issues resident in data stream applications, this paper added the element of age to sliding window data. Our contributions include two incremental quantization techniques, IVQ and DBQ . These linear heuristics, while more desirable due to the real-time constraints of most stream applications, have been shown to suffer only minor losses in data representation accuracy compared to the optimal solution. Finally, we described the

Ladder Approach to apply the correct ages to specific portions of data, allowing our system to maintain a more robust sliding window.

The methods discussed in this paper have largely relied on the assumption that data arrival rates are constant. But as we well know, this situation is hardly typical, and the problems they cause should not be overlooked. For instance, a sharp increase in data arrival rates will saturate the sliding window with all young elements, while older elements will already have been dropped from the synopsis. Depending on the application, a sliding window that hypothetically only contains elements arriving in the past k milliseconds is probably not all that useful. We believe that our aging method can also be applied here for supporting real-time based sliding windows under dynamic environments. Future work on this topic will certainly be under this direction.

References

- [1] F. Altıparmak, D. Chiu, and H. Ferhatosmanoglu. Incremental quantization for aging data streams. Technical Report OSU-CISRC-3/07-TR20, The Ohio State University.
- [2] B. Babcock, S. Babu, M. Datar, R. Motwani, and J. Widom. Models and issues in data stream systems. In *PODS*, pages 1–16, New York, NY, USA, 2002. ACM Press.
- [3] B. Babcock, M. Datar, and R. Motwani. Sampling from a moving window over streaming data. In *SODA*, 2002.
- [4] B. Babcock, M. Datar, and R. Motwani. Load shedding for aggregation queries over data streams. In *ICDE*, page 350, Washington, DC, USA, 2004.
- [5] L. Chen and G. Agrawal. Supporting self-adaptation in streaming data mining applications. In *IPDPS 2006*. IEEE.
- [6] S. Guha. Space efficiency in synopsis construction algorithms. In *VLDB*, pages 409–420, 2005.
- [7] S. Guha, N. Koudas, and K. Shim. Data-streams and histograms. In *STOC*, pages 471–475, New York, NY, USA, 2001. ACM Press.
- [8] P. Indyk, N. Koudas, and S. Muthukrishnan. Identifying representative trends in massive time series data sets using sketches. In *VLDB*, September, 2000.
- [9] Y. E. Ioannidis and V. Poosala. Balancing histogram optimality and practicality for query result size estimation. In *SIGMOD*, pages 233–244, New York, NY, USA, 1995. ACM Press.
- [10] H. V. Jagadish, N. Koudas, S. Muthukrishnan, V. Poosala, K. C. Sevcik, and T. Suel. Optimal histograms with quality guarantees. In *VLDB*, pages 275–286, 24–27 1998.
- [11] S. P. Lloyd. Least squares quantization in pcm. *IEEE Transactions on Information Theory*, 28:127–135, Mar. 1982.
- [12] Y. Ogras and H. Ferhatosmanoglu. Online summarization of dynamic time series data. *The VLDB Journal*, 15(1):84–98, 2006.
- [13] N. Tatbul, U. Cetintemel, S. Zdonik, M. Chemiack, and M. Stonebraker. Load shedding in a data stream manager. In *VLDB*, Berlin, Germany, 2003.